

Balanced Binary Tree Crate Report

Mackenzie Malainey - `mmalaine@ualberta.ca`

Lora Ma - `lora@ualberta.ca`

Benjamin Kong - `bkong@ualberta.ca`

March 7, 2022

Contents

1	Major Innovations	3
1.1	Custom TreeBalance Implementations	3
2	Design	4
2.1	Rationale	4
2.2	Analysis	5
2.2.1	Crate Structure	5
2.2.2	Tree Structure	6
2.3	Extension	7
3	Known Issues	8
3.1	TreeBalance Trait Implementation	8
4	Benchmarking	9
4.1	Red Black Tree	9
4.2	AVL Tree	12
4.3	Comparison	14
5	User Manual	16
5.1	Tree	16
5.2	Pre-Defined Types	16
5.3	TreeBalance Trait Implementation	17
5.4	CLI Demo	17

1 Major Innovations

1.1 Custom TreeBalance Implementations

As a result of working towards our goal of implementing a Red Black Tree and AVL Tree with shared components, we designed a system that allows library users to implement their own algorithm to perform adjustments to the tree structure after insertion or deletion. This system allows users to implement the `TreeBalance` trait within their own code and use that for performing balances on the tree after the structure changes. While the library does already expose the common Red Black Tree and AVL Tree types, because of how they were implemented we were able to expose this method to create a “custom balance” so that users may implement ones that better suit their needs.

For example: Instead of a balanced tree, it might be better for the user to have a binary search tree that is skewed to the right. A case for this could be when working a very large tree that will rarely search for elements that would be on the far right, but is mostly going to be used to search for elements on the far left. Providing a custom balancing algorithm to skew the tree such that the far left is closer to the middle could provide a performance boost in the long term. The exposure of the `TreeBalance` trait allows library users to do just that, without having to be concerned about implementing the entire tree. Furthermore, they still have access to the AVL Tree and Red Black Tree types if those end up being required as well.

Additionally, from the `TreeBalance` trait we implemented a non-balancing tree (plain binary search tree) and exposed that type as well.

2 Design

2.1 Rationale

This crate was designed based on the requirement of the implementation of both a Red Black Tree and an AVL Tree for generalized types. In essence, a Red Black Tree and an AVL Tree are both Binary Search Trees but with balancing checks and operations that are executed whenever the structure of the tree changes (an insertion or deletion occurs). From this the idea of a balanced tree can be broke into two components: a Binary Search Tree and a balancing algorithm.

Further investigation into the nature of the AVL and Red Black algorithms displayed that the two balancing algorithms both perform very similar operations:

- Inspect a subtree that was affected by the operation
- Perform a rotation on the first/second levels of the subtree

This revealed that neither algorithm was proactive in tree operations but instead are reactive (has no say in how the actual is operation performed, instead uses the changes in the tree structure to rebalance after the fact). From that, it became clear that performing operations on the tree and the act of rebalancing do not influence each other except through their changes to the tree. Therefore tree operations can be opaque to the balancing algorithms (aside from performing rotations).

Additionally, we find that as long as the balancing algorithms are provided a way to inspect the structure of the tree they do not require access to any of the actual structure of it, aside from a method to perform rotations on the tree. This became the motivation for the `tree::inspect` components that hide away the tree's internal code structure so that the balancing algorithm is not able to break the contract of a binary search tree. If the internal workings of the tree are safely abstracted away then any balancing algorithm can

be implemented with the guarantee that the tree will remain a binary search tree after rebalancing. This allowed us to not only safely implement the AVL Tree and Red Black Tree algorithms with those guarantees, but expose the `TreeBalance` trait to the public for custom implementation. And any operations that could break the tree's structure are guarded using panics (not recoverable errors as contract breaking changes are always preventable)

However, there was one decision still to be identified: how would we associate a balancing function with a given tree. An AVL Tree only requires information on a subtree's structure that is already defined (its' height) and which subtree to start looking at. Opposite to this, a Red Black Tree requires metadata to be associated with each node that the algorithm must have read/write access to. From this we realized that a balancing algorithm should be allowed access to a mutable metadata stored at each node if it requires it.

To summarize our decision was to create a binary search tree whose programmatical structure and how it performed operations were opaque to any balancing algorithm, but would use an associated balancing algorithm (defined at compile-time and whose implementation is mostly opaque to the tree) after a successful tree operation to perform rebalancing on the tree (with optional mutable access to metadata only used and adjusted by the balancer). We also required error handling on public `Tree` methods since cases where a delete or insertion fail can occur acceptably.

2.2 Analysis

2.2.1 Crate Structure

There are four main things this crate exposes:

- `Tree` - the public interface for a binary search tree
- `AVLTree` - implementation of an AVLTree on a `Tree` using the `TreeBalance` trait

- `RedBlackTree` - implementation of an `RedBlackTree` on a `Tree` using the `TreeBalance` trait
- `TreeBalance` - trait definition for implementing a type a `Tree` will use to perform balancing operations

The crate is organized into the following modules:

- `lib` - top level library that exposes necessary components outside the crate
- `main` - CLI Demo code
- `avl` - the AVL Tree implementation of the `TreeBalance` trait
- `redblack` - the Red Black Tree implementation of the `TreeBalance` trait
- `tree` - `Tree` definition and implementation, submodules are further implementation dependencies
- `tree::inspect` - holds trait definitions, structs, and enums required and related to the `TreeBalance` trait
- `tree::ops` - holds functions to perform operations on a tree
- `tree::node` - holds struct definition for a node in the tree

2.2.2 Tree Structure

A `Tree` acts as both the public facing interface for working with a binary search tree as well as acting as a way a container for pointers to a node's child objects as it will possibly contain a pointer to a `TreeNode`. A `TreeNode` is an actual node in a tree stored on the heap that contains the key, parent and children pointers, and metadata regarding the tree's structure. Any type that implements the `TreeBalance` trait can be used associated with

a `Tree` type so it is used to perform balancing after operations. As a method to safeguard a `Tree`'s internal structure when a rebalance is requested a `NodeInspector` is passed into the relevant `TreeBalance` trait function as a handle that can be used to inspect the tree's structure and perform rotations as necessary. A `Tree` must have an associated `TreeBalance` trait with it (types are associated at compiled time through generics).

2.3 Extension

As mentioned already the code is fairly extensible by custom implementing the `TreeBalance` trait and using the associated `NodeInspector` and related types. However, in regards to adjustments required to implement a more complex type of tree such, as a 2-3-4 Tree or a B-Tree, we would need to significantly change the design of the `Tree` structure. First of we would need to generalize how child nodes are indexed, therefore allowing more than two children per node (and it becoming opaque to the operation functions, to be implemented outside of the structure). Secondly, nodes will need to be multi-keyed meaning we will need to adjust how we handle the `TreeNode` to allow for multiple keys. Lastly, the `Tree` structure invokes the balancing algorithm in as a reaction to structure change, a B-Tree and 2-3-4 Tree are balanced proactively (before the insertion is made). This would require providing the balancing algorithm access during search now on top of when travelling back up the tree (if we want to continue to provide support for the AVLTree case). Furthermore we would also need to expose a splitting operation for multi-key nodes. Overall, for complex multikey and multi-path nodes this is not easily extensible, a large rewrite of the library would be required. Although this could be made easier by instead of generalizing the node type to handle all cases, we generalize the input and delete operations to handle multiple different node types (binary tree node, multi-key node, etc) and have multiple node types that handle invoking the

3 Known Issues

3.1 TreeBalance Trait Implementation

The implementation of the `TreeBalance` trait is not as robust as it could be and is lacking some important features. However these features have no impact on whether or not the AVL Tree and Red Black Tree work correctly.

The issue is that simply right now you cannot:

- Travel more than one node (other than going to root) when returning the next position to go on the tree from a rebalance function.
- Travel to a child node when returning the next position to go on the tree from the `rebalance_insert` function.


This feature loss might cause a minor performance impact for the Red Black Tree as it has a case where it is beneficial to move up two parents at once. However we have implemented a fix for this, but it might be slower as it still requires checking the surrounding node's colors before deciding it is that special case and moves up a parent again.

4 Benchmarking

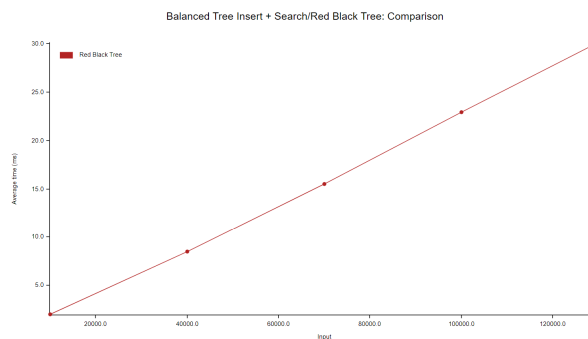
These stats are generated on a Windows 10 PC with an AMD Ryzen 7 5800X CPU and 32GB of RAM.

4.1 Red Black Tree

The following chart shows the relationship between function/parameter and iteration time. The thickness of the shaded region indicates the probability that a measurement of the given function/parameter would take a particular length of time.





The following chart shows the mean measured time for each function as the input (or the size of the input) increases.



Balanced Tree Insert + Search/Red Black Tree/10000

Balanced Tree Insert + Search/Red Black Tree/10000





Additional Statistics:

	Lower bound	Estimate	Upper bound
R ²	0.0413393	0.0429060	0.0413178
Mean	2.0363 ms	2.0484 ms	2.0606 ms
Std. Dev.	53.514 us	61.029 us	67.647 us
Median	2.0386 ms	2.0564 ms	2.0666 ms
MAD	42.045 us	57.198 us	79.834 us

Balanced Tree Insert + Search/Red Black Tree/40000

Balanced Tree Insert + Search/Red Black Tree/40000





Additional Statistics:

	Lower bound	Estimate	Upper bound
R ²	0.0175126	0.0181166	0.0173994
Mean	8.4820 ms	8.5108 ms	8.5423 ms
Std. Dev.	116.27 us	154.43 us	188.29 us
Median	8.4497 ms	8.4793 ms	8.5088 ms
MAD	89.742 us	116.58 us	141.31 us

Balanced Tree Insert + Search/Red Black Tree/70000

Balanced Tree Insert + Search/Red Black Tree/70000





Additional Statistics:

	Lower bound	Estimate	Upper bound
R ²	0.0079850	0.0082677	0.0079416
Mean	15.441 ms	15.483 ms	15.529 ms
Std. Dev.	181.34 us	223.80 us	262.36 us
Median	15.399 ms	15.447 ms	15.492 ms
MAD	157.48 us	208.43 us	238.97 us

Balanced Tree Insert + Search/Red Black Tree/100000

Balanced Tree Insert + Search/Red Black Tree/100000





Additional Statistics:

	Lower bound	Estimate	Upper bound
R^2	0.0023715	0.0024480	0.0023469
Mean	22.779 ms	22.920 ms	23.083 ms
Std. Dev.	481.65 us	788.17 us	1.0770 ms
Median	22.683 ms	22.745 ms	22.813 ms
MAD	268.76 us	335.29 us	453.01 us

Balanced Tree Insert + Search/Red Black Tree/130000

Balanced Tree Insert + Search/Red Black Tree/130000




Additional Statistics:

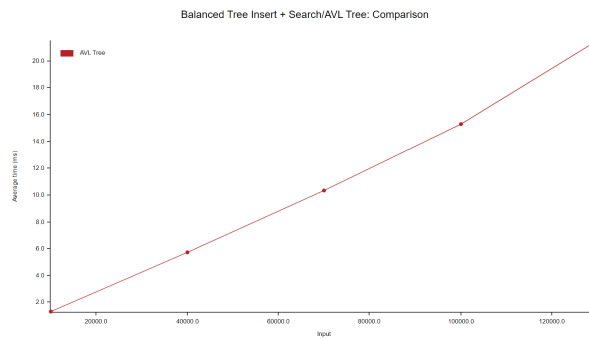
	Lower bound	Estimate	Upper bound
R^2	0.0004869	0.0005027	0.0004811
Mean	30.135 ms	30.244 ms	30.371 ms
Std. Dev.	387.14 us	605.00 us	847.98 us
Median	30.011 ms	30.097 ms	30.181 ms
MAD	281.43 us	382.29 us	510.87 us

4.2 AVL Tree

The following chart shows the relationship between function/parameter and iteration time. The thickness of the shaded region indicates the probability that a measurement of the given function/parameter would take a particular length of time.





The following chart shows the mean measured time for each function as the input (or the size of the input) increases.



Balanced Tree Insert + Search/AVL Tree/10000

Balanced Tree Insert + Search/AVL Tree/10000





Additional Statistics:

	Lower bound	Estimate	Upper bound
Slope	1.3182 ms	1.3222 ms	1.3261 ms
R ²	0.9758104	0.9770795	0.9758665
Mean	1.3142 ms	1.3219 ms	1.3332 ms
Std. Dev.	18.227 us	50.010 us	80.670 us
Median	1.3123 ms	1.3164 ms	1.3199 ms
MAD	10.378 us	13.488 us	21.579 us

Balanced Tree Insert + Search/AVL Tree/40000

Balanced Tree Insert + Search/AVL Tree/40000





Additional Statistics:

	Lower bound	Estimate	Upper bound
R ²	0.0138371	0.0143212	0.0137619
Mean	5.7015 ms	5.7156 ms	5.7309 ms
Std. Dev.	55.451 us	75.388 us	94.965 us
Median	5.6988 ms	5.7085 ms	5.7210 ms
MAD	43.504 us	56.646 us	74.119 us

Balanced Tree Insert + Search/AVL Tree/70000

Balanced Tree Insert + Search/AVL Tree/70000





Additional Statistics:

	Lower bound	Estimate	Upper bound
R ²	0.0575267	0.0595013	0.0573613
Mean	10.280 ms	10.309 ms	10.339 ms
Std. Dev.	125.65 us	150.65 us	171.90 us
Median	10.250 ms	10.267 ms	10.307 ms
MAD	89.012 us	116.40 us	151.54 us

Balanced Tree Insert + Search/AVL Tree/100000

Balanced Tree Insert + Search/AVL Tree/100000





Additional Statistics:

	Lower bound	Estimate	Upper bound
R ²	0.0003327	0.0003426	0.0003275
Mean	15.192 ms	15.260 ms	15.345 ms
Std. Dev.	195.36 us	398.57 us	590.01 us
Median	15.145 ms	15.172 ms	15.205 ms
MAD	117.74 us	178.06 us	221.97 us

Balanced Tree Insert + Search/AVL Tree/130000

Balanced Tree Insert + Search/AVL Tree/130000



Additional Statistics:

	Lower bound	Estimate	Upper bound
R ²	0.0014983	0.0015482	0.0014848
Mean	20.779 ms	21.600 ms	22.531 ms
Std. Dev.	3.0444 ms	4.5238 ms	5.6385 ms
Median	19.876 ms	20.001 ms	20.137 ms
MAD	320.58 us	459.20 us	737.63 us

4.3 Comparison

For every case the AVL Tree performs noticeably better than the Red Black Tree illustrating that the AVL Tree is more efficient. However, this conclusion really can only be said for this worst possible “insert and search” case. This does not provide insight on best case behavior or on random behavior. Furthermore we have no benchmarks for any deletion case. These additional benchmarks would be valuable data for getting a more complete picture on how the two data structures differ performance wise. As well, if we have any knowledge of the data we expect to receive we can use the additional bookmarks to select the best data structure for the case we are most likely to match.

Using a binary search tree as a baseline measurement is not likely to be necessary. We

know that the worst case complexity of the AVL Tree and Red Black Tree are better than that of a binary search tree and the comparison is between the two trees. However if we include more benchmarks for other cases having a baseline might be more beneficial so we may get a better picture of how more general cases actually compare. It might also be beneficial to include benchmarks for 2-3-4 Trees and B-Trees. Again, this way we can get more information regarding what sort of tree would be the best fit for our use case (provided we know anything in advance about how the tree will get used most commonly).

5 User Manual

5.1 Tree

The `Tree T, U` type represents a binary search tree to search `T` and is balanced by `U`. The following are methods of the public interface for manipulating a `Tree` instance:

- `Tree::new` - constructs an empty tree
- `Tree::is_empty` - returns true if the tree has no nodes
- `Tree::insert` - inserts a key, rebalancing using `U`
- `Tree::search` - returns true if the key was found in the tree
- `Tree::delete` - removes and returns a key, rebalancing using `U`
- `Tree::height` - returns the height of the tree
- `Tree::leaves` - returns the number of leaf nodes in the tree

5.2 Pre-Defined Types

The following are other predefined types for basic `Tree` usage

- `avl::AVLBalance` - `TreeBalance` type using AVL Tree rules
- `AVLTree` - typedef for tree that uses `AVLBalance`
- `redblack::RedBlackTree` - `TreeBalance` type using Red Black Tree rules
- `RedBlackTree` - typedef for tree that uses `RedBlackBalance`
- `BinarySearchTree` - typedef for a tree that doesn't perform any balancing

5.3 TreeBalance Trait Implementation

The `TreeBalance` trait can be used to perform rebalancing on a tree using custom rules. A type that implements the trait and is associated with a `Tree` instance will have the trait functions called to perform rebalancing. Review `TreeBalance` and relevant type documentation for more information as we will not describe it further here.

5.4 CLI Demo

We have built a CLI that allows users to test both the Red Black Tree and the AVL Tree. Upon starting the program (located in `main.rs`), the user will be greeted by a welcome message. The user will also be prompted to select which type of tree they would like to test. In order to test a Red Black Tree, the user needs to enter '1'; to test an AVL Tree, the user needs to enter '2'.

Upon selecting a tree, the user will be presented with a list of commands. Here is a summary of the commands available:

- `insert <n>` - insert `<n>` into the tree
- `delete <n>` - delete `<n>` from the tree
- `search <n>` - output whether `<n>` is in the tree or not
- `clear` - clear all elements from the tree
- `isempty` - output whether the tree is empty or not
- `height` - output the height of the tree
- `leaves` - output the number of leaves in the tree
- `print tree` - print formatted tree

- `print inorder` - print elements using inorder traversal
- `switch` - select a new tree type
- `help` - reprint this menu
- `exit` - exit program

For example, let's say we want to select an AVL Tree and insert 3, 5, and 2 into the tree. We would first enter '2' followed by `insert 3`, `insert 5`, and `insert 2`.